


What's in a GPU?

- Heterogeneous chip multi-processor
- Tuned for graphics
- Fast computing




What's in a GPU?

- History :
 - 1990's real time 3D rendering
 - Video games, Movies, etc.
 - Computationally VERY EXPENSIVE!
 - Before 1990's graphics were done on CPU's as well as the framerate!
 - Reduced quality images
 - No 3D rendering

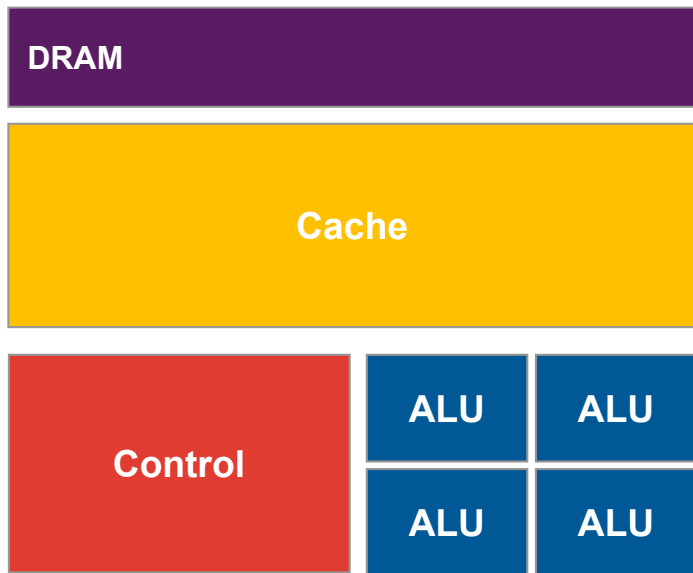
What's in a GPU?

- History :
 - 1990's real time 3D rendering
 - Before 1990's graphics were done on CPU's




What's in a GPU?

CPU



GPU




Drug discovery with MegaDOCK on GPUs

Drug discovery with MegaDOCK on GPUs

- **GOAL**

- Use hip to port a real application to AMD GPUs
 - Evaluate performance due to portability
 - General portability issues
- Next generation of GPU architectures – Frontier ORNL



Drug discovery with MegaDOCK on GPUs

- **GOAL**



- Preparing for the next generation of GPU architectures
 - Frontier ORNL in 2021
- Effort of porting applications for these systems
- Applications working with maximum performance
- Broad range of applications



Drug discovery with MegaDOCK on GPUs

- **OUTLINE**


- What is molecular docking?
- What is MEGADOCK?
- Algorithm
- Why to use hip?
- Timings
- Zlab benchmark



Drug discovery with MegaDOCK on GPUs

- **WHAT IS MOLECULAR DOCKING?**

- Computational tool used to predict protein interactions
 - Determine how two proteins interact in their native environment
 - Aids to understand biological processes at the molecular level
 - Helps to design drugs to fight certain diseases
 - Important to predict conformation and protein interactions
 - COVID-19




Drug discovery with MegaDOCK on GPUs

- **WHAT IS MOLECULAR DOCKING?**




LIGAND

+



RECEPTOR

=



Drug discovery with MegaDOCK on GPUs

- **APPLICATIONS**

- Different types of docking:
 - Rigid
 - Semi-rigid
 - Flexible

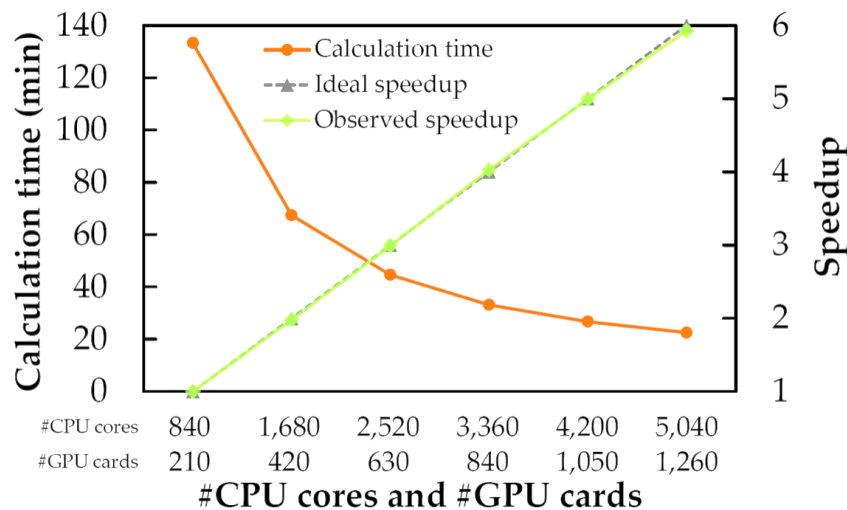
} Different computation demands
- Within the different applications:
 - **MEGADOCK**
 - Fast
 - Optimized to the GPU by NVIDIA
 - Not very difficult to port

Drug discovery with MegaDOCK on GPUs

- **MEGADOCK**

- Single node CPU & GPU
- Multi node CPU and GPU


- **SCALABILITY**



- MPI/OpenMP & CUDA
- Linear speed up
- Matches the ideal curve
- Speed up of 29x with the GPU usage

Drug discovery with MegaDOCK on GPUs

- ALGORITHM



calcrg.cpp	control.cpp	cpu_time.h	docking.cpp	fft_process.cu	ligand.cpp	main.cu	parallel.h	pdb_entry.cpp	protein.h
constant.h	control.h	cuda_kernel.cu	docking.h	fft_process.h	ligand.h	mpidp.cpp	parameter.cpp	pdb_entry.h	receptor.cpp
contact.py	cpu_time.cpp	decoygen.cpp	fft_process.cpp	helper_cuda.h	main.cpp	mpidp.h	parameter.h	protein.cpp	receptor.h




cuda_kernel.cu fft_process.cu main.cu

Drug discovery with MegaDOCK on GPUs

- **IMPLEMENTATION**

- Highly optimized for **Nvidia GPUs using CUDA**



Drug discovery with MegaDOCK on GPUs

- **WHY TO USE HIP?**

- **Example HIP:**

- **main.cu**

Original Cuda

```
if(nogpu_flag != 1) {
    checkCudaErrors( cudaGetDeviceCount(&device_count_gpu) );
    if (device_count_gpu == 0) {
        fprintf(stderr, "GPU Error: no devices supporting CUDA.\n");
        exit(-1);
    }

    cudaDeviceProp deviceProp;
    checkCudaErrors( cudaGetDeviceProperties(&deviceProp, 0));
    if (deviceProp.major < 1) {
        fprintf(stderr, "GPU Error: device does not support CUDA.\n");
        exit(-1);
    }

    cudaSetDeviceFlags(cudaDeviceMapHost);
    fprintf(stdout, "# Using CUDA device %d: %s\n", 0, deviceProp.name);
    cudaSetDevice(0);
    //fprintf(stdout, "# Init CUDA device OK.\n");

    int cufft_version;
    cufftGetVersion(&cufft_version);
    printf("# CUFFT version : %d\n", cufft_version);
}
#endif
```

15

After hipify

```
if(nogpu_flag != 1) {
    checkCudaErrors( hipGetDeviceCount(&device_count_gpu) );
    if (device_count_gpu == 0) {
        fprintf(stderr, "GPU Error: no devices supporting CUDA.\n");
        exit(-1);
    }

    hipDeviceProp_t deviceProp;
    checkCudaErrors( hipGetDeviceProperties(&deviceProp, 0));
    if (deviceProp.major < 1) {
        fprintf(stderr, "GPU Error: device does not support CUDA.\n");
        exit(-1);
    }

    hipSetDeviceFlags(hipDeviceMapHost);
    fprintf(stdout, "# Using CUDA device %d: %s\n", 0, deviceProp.name);
    hipSetDevice(0);
    //fprintf(stdout, "# Init CUDA device OK.\n");

    int cufft_version;
    hipfftGetVersion(&cufft_version);
    printf("# CUFFT version : %d\n", cufft_version);
}
#endif
```



UNIVERSITY OF AMSTERDAM



Drug discovery with MegaDOCK on GPUs

- TIMINGS FOR RIGID DOCKING BENCHMARK**

	¹ Summit-dev	Kleurplaat (AMD-ROME)
Single CPU	89.24s	23.96s
Single (CPU & GPU)	5.79s	5.39s

- No loss of performance with Hipify
- Significant speed up from CPU to GPU (~4 times for AMD-ROME)

Drug discovery with MegaDOCK on GPUs

- ZLAB DOCKING BENCHMARK**

	CPU	GPU
RIGID	138.75s	8.62s
SEMI-RIGID	337.42s	12.07s
FLEXIBLE	607.88s	17.94s

AMD

Drug discovery with MegaDOCK on GPUs

- **CONCLUSIONS**

- No loss of performance by using HIP
- Significant speed up by using the AMD GPUs
- Easy to use
- Solves issues with portability
- It is worth trying

AMD

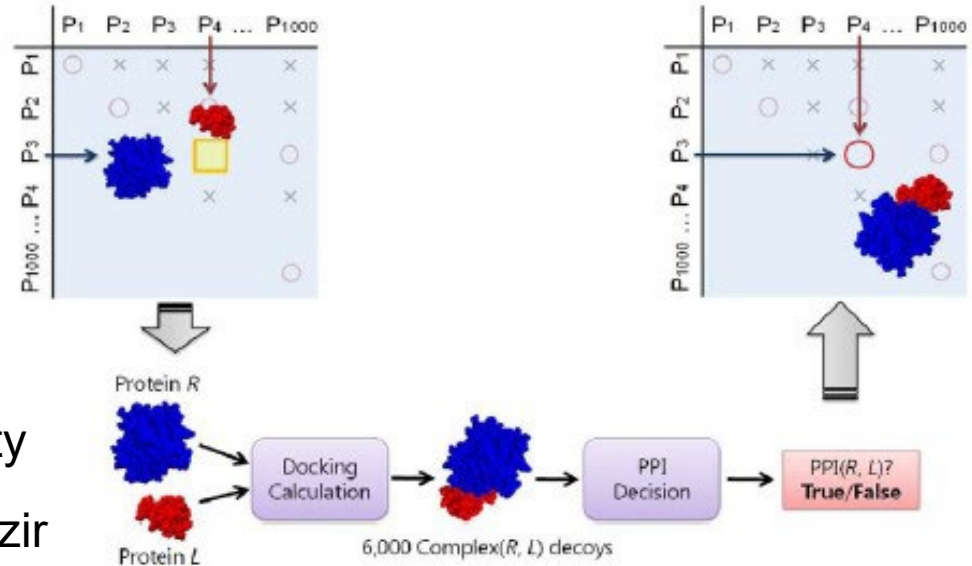
THANK YOU FOR YOUR TIME

AMD

Porting applications to the GPU using Hip

- Algorithm

- Receptor (R)
- Ligand (L)
- FFT based Docking
- Rigid
- Pairwise Shape complementarity
- Scoring Method: Katchalski-Katzir



Porting applications to the GPU using Hip

- **Algorithm**
 - Scoring
- Three-dimensional voxels
- Square & Circle : area occupied by proteins
- Voxels with only the squares represent unoccupied space

